KIMI4 PRACTICE EXAM

ANSWER KEY WITH EXPLANATIONS

Single-choice items score 1 point. [SELECT TWO/THREE] items score only if every required option is exactly right - no partial credit, like the real exam. Objective references map each question to the official exam outline.

QUICK MARKING TABLE

Q1: A	Q2: C	Q3: B	Q4: D	Q5: B	Q6: A, B
Q7: C	Q8: A	Q9: D	Q10: B	Q11: A	Q12: B, D
Q13: C	Q14: D	Q15: C	Q16: B	Q17: A	Q18: C
Q19: B	Q20: D	Q21: B	Q22: C	Q23: A	Q24: A, B
Q25: D	Q26: B	Q27: C	Q28: D	Q29: A	Q30: B
Q31: D	Q32: B	Q33: B	Q34: D	Q35: A	Q36: C
Q37: B	Q38: A	Q39: A	Q40: C	Q41: B	Q42: C
Q43: D	Q44: A	Q45: B, D	Q46: B	Q47: D	Q48: A, C
Q49: D	Q50: A	Q51: C	Q52: B	Q53: A	Q54: D
Q55: B	Q56: C	Q57: A, B, C	Q58: D	Q59: A	Q60: A

EXPLANATIONS - DOMAIN 1 - CONNECT TO AND TRANSFORM DATA (Q1-Q14, 24%)

Q1 - Correct: A

Objective: 1.1 - Connect to files

JSON connections open with the Select Schema Levels dialog; the levels you tick are flattened into tables which appear in the left pane. Choosing the right level controls the grain of the resulting table.

Option D confuses JSON with Custom SQL (a relational-database feature); JSON needs no conversion and is fully supported in Desktop.

Q2 - Correct: C

Objective: 1.1 - Live vs extract / packaging

A .twb is XML only - with a live connection it points at a database the consultant cannot reach. Packaging as .twbx bundles the workbook with its local files and extracts, so it opens offline. This is the classic portability reason for extracts.

Related but different file: .tdsx packages a DATA SOURCE, not a workbook.

Q3 - Correct: B

Objective: 1.1 - Live vs extract

An extract decouples the dashboards from warehouse availability (works during the outage) and speed (Hyper is fast during the day), and a nightly refresh satisfies the data-as-of-yesterday requirement.

Extracts win on performance, portability and availability; live wins when real-time data is mandatory.

Q4 - Correct: D

Objective: 1.2 - Joins

The physical join dialog lets each clause use =, <>, <, >, <= or >= (and you can stack multiple clauses). That is exactly how 'event falls inside interval' problems are modeled: Product = Product AND Order Date >= Start Date AND Order Date <= End Date.

Relationships/blends don't change this - blend matches on shared members at aggregate level only.

Q5 - Correct: B

Objective: 1.2 - Blending

A blend relationship only works in a view when the linking field is ACTIVATED - the chain icon beside the secondary field in the Data pane. Auto-detected links can be inactive; clicking the icon toggles the link.

If the field NAMES differed you would use Data menu > Edit Blend Relationships instead. A repeated grand total is the signature of an unlinked blend dimension.

Q6 [SELECT TWO] - Correct: A, B

Objective: 1.2 - Data quality

Profiling is about completeness (null share) and consistency/accuracy (odd member counts, duplicates of the same entity under different spellings). Both are visible in the profile panes of Tableau Prep and in the data grid.

Response time and permissions are governance/performance topics, not data quality.

Q7 - Correct: C

Objective: 1.3 - Prep cleaning

Data roles in Prep (Email, URL, geographic roles, etc.) validate a field's values against a known format and mark the invalid ones in the profile pane - the fastest route to finding 'impossible' values.

Group Values handles near-duplicate spellings; data roles handle format validity.

Q8 - Correct: A

Objective: 1.3 - Prep cleaning

Prep's one-click cleaning operations include Trim Spaces, Remove Letters, Remove Numbers, Remove Punctuation and Remove All Spaces - each edits values in place without writing a calculation.

Choose Remove Punctuation here; Remove Numbers would have done the opposite of what is needed.

Q9 - Correct: D

Objective: 1.2 / 1.4 - Field organization

The Data pane groups by data-source table by default; right-clicking the pane's background (or the dimensions area) lets you group by Folder instead, after which 'Create Folder' organizes fields into your own named groups.

Hierarchies are for drill-down, sets for row membership, and combined fields concatenate keys - none organizes the pane itself.

Q10 - Correct: B

Objective: 1.4 - Default field properties

Default Properties > Sort stores the sort rule with the field, so every view inherits it. Manual sort lets you drag the four regions into the exact business order.

Aliases change display text, not order; per-view sorts must be redone on each sheet.

Q11 - Correct: A

Objective: 1.3 - Prep joins

The Join step's workspace shows the join type diagram plus a summary of the join clause results, including counts of matched and excluded values, so you can validate and tune the join before running the flow.

The Changes pane records edits; it never shows row-match statistics.

Q12 [SELECT TWO] - Correct: B, D

Objective: 1.2 - Relationships vs joins

A physical join is computed up front into a single denormalized table; relationships stay two logical tables combined per-query at the right grain. That is why one-to-many joins inflate 'one-side' measures while relationships do not.

The exam's standing advice: default to relationships; join physically only when you truly need the fixed row-level result.

Q13 - Correct: C

Objective: 1.1 - Connect to extracts

An extract is itself a data source: connect to it directly with the file picker (the Hyper engine reads it at full speed). No flow, server, or conversion is needed.

The same dialog also opens .tde, .csv, .xlsx, .json, .pdf and statistical files.

Q14 - Correct: D

Objective: 1.1 - Connect to files

A text file's separator, text qualifier, character set and locale are properties of the connection (right-click the file on the data source page > Text File Properties). Setting the semicolon separator and the correct locale makes the file read cleanly.

Data Interpreter removes banners/footers; it does not change delimiter or decimal conventions.

EXPLANATIONS - DOMAIN 2 - EXPLORE AND ANALYZE DATA (Q15-Q39, 41%)

Q15 - Correct: C

Objective: 2.2/2.3 - Quick table calculations

YTD Total is a running total PARTITIONED by year: Jan 5 -> Feb 15 -> Mar 30, and January 2025 starts a new partition, showing just that January's value. Knowing that quick table calcs auto-partition by the date hierarchy is the trick.

30 is the running sum of 5+10+15; 'forever accumulating' would be plain Running Total.

Q16 - Correct: B

Objective: 2.2 - Difference quick table calc

Difference From computes current minus previous: 30 - 50 = -20. January shows NULL because it has no previous mark.

The -40% answer is the PERCENT Difference variant (-20/50); mixing the two up is the classic trap on this family of questions.

Q17 - Correct: A

Objective: 2.3 - Window functions

Offsets -2..0 mean 'the two marks before me plus me': Jan 10 + Feb 20 + Mar 30 = 60. Offset 0 is always the current mark; negative offsets look backward in the addressing direction.

90 would require -3..0 (April would give 90 for -2..0 - wait, April would give 20+30+40=90); 30 forgets the window entirely.

Q18 - Correct: C

Objective: 2.3 - LOOKUP

LOOKUP(value, +1) reads the NEXT mark in the partition. For the final mark there is no next mark, so the result is NULL. Positive offsets look forward; negative look back.

No error is raised - out-of-range lookups simply return NULL, which is exactly why LOOKUP is safe at edge marks.

Q19 - Correct: B

Objective: 2.3 - Running functions

RUNNING_MAX carries the largest value seen so far: 30 -> 30 -> 50 -> 50. The fourth mark therefore shows 50, not its own 20.

This is the function behind 'high-water mark' views. RUNNING_MIN would mirror it with the lowest value.

Q20 - Correct: D

Objective: 2.2/2.3 - Compute Using direction

Addressing sets the direction in which the percentage accumulates: computing DOWN the Region rows makes each month column its own partition summing to 100%. 'Table (across)' would make each REGION ROW sum to 100% instead.

Cell gives 100% everywhere and Pane splits at pane boundaries - the exam's favorite wrong answers for direction questions.

Q21 - Correct: B

Objective: 2.8 - INCLUDE LOD

INCLUDE starts from the viz grain and ZOOMS FINER. Region is a COARSER grain than State, so it adds nothing: the per-(state, region) value equals the state value. This is the single most misunderstood INCLUDE fact.

To get the region total beside each state you need REMOVE semantics: {EXCLUDE [State] : SUM([Sales])} or {FIXED [Region] : SUM([Sales])}.

Q22 - Correct: C

Objective: 2.8 - EXCLUDE LOD + context

Context filters sit ABOVE every LOD stage (fixed and scoped alike). With Boston in context, the EXCLUDE [City] computation runs over the context-filtered rows only: 100. The 'EXCLUDE ignores context' instinct is the trap - only DIMENSION and MEASURE filters are ignored by EXCLUDE/INCLUDE.

Contrast: {FIXED ...} ignores dimension filters but still obeys context, data-source and extract filters.

Q23 - Correct: A

Objective: 2.1 + 2.8 - LOD flag calculation

The inner LOD freezes each customer's first date; wrapping both sides in YEAR makes it a cohort-year comparison. Because the LOD is FIXED, the flag stays stable whatever dimensions the view later adds.

Option B flags the customers' latest year; option C compares against the view's own minimum; option D misuses a table-calc helper as data.

Q24 [SELECT TWO] - Correct: A, B

Objective: 2.8 - FIXED as dimension

A FIXED LOD whose inner aggregate yields one non-measure value per member can become a dimension - the classic cohort patterns A and B. That is what makes cohort charts possible at all.

Option C produces a measure (grand total), and option D is a table calculation, which can never be a dimension.

Q25 - Correct: D

Objective: 2.1 - Aggregates

PERCENTILE(expression, number) is the built-in aggregate (number between 0 and 1); 0.95 returns the 95th percentile of the underlying rows. There is no QUANTILE() function in Tableau, and scaling MAX or weighting MEDIAN are not percentile formulas.

Note PERCENTILE is an aggregate like AVG - for live connections to some databases its exact computation can vary; on extracts it is always available.

Q26 - Correct: B

Objective: 2.1 - Predictive functions

MODEL_QUANTILE and MODEL_PERCENTILE are Tableau's predictive modeling functions; they run as table calculations over the data in the view (quantile: predicted value at a probability; percentile: probability of a value). SCRIPT_* functions push work out to R/Python via an external service.

Knowing the pair's names is enough for the exam - deep syntax is not tested.

Q27 - Correct: C

Objective: 2.4 - Filters

Relative-date filters keep a moving window anchored to today (or to a chosen anchor date): 'last N days', 'this quarter', 'year to date', etc. The view re-evaluates on every load.

Condition/wildcard operate on members and would hard-code dates; Top N has nothing to do with time windows.

Q28 - Correct: D

Objective: 2.4 - Measure filters

Measure filters offer 'At least', 'At most' and a 'Range of values' with both endpoints, applied to the aggregated value per mark. One range filter covers the whole 'between' requirement.

Bins bucket continuous values for histograms - they filter nothing.

Q29 - Correct: A

Objective: 2.4 - Filters across sheets

Apply to Worksheets offers Only This Worksheet / Selected Worksheets / All Using This Data Source. 'All Using This Data Source' makes the single filter control every sheet on that source automatically - including sheets built later.

Use 'Selected Worksheets' when only a subset of sheets should respond.

Q30 - Correct: B

Objective: 2.4 - Top / Bottom N

The Top tab toggles between Top N and Bottom N and lets you pick the measure + aggregation the ranking is based on (here AVG(Profit)). It recomputes as data changes - manual hiding does not.

A conditional filter or a dynamic set could imitate it, but the exam's 'designed-for-it' answer is the Top tab.

Q31 - Correct: D

Objective: 2.4 - Conditional filters

The Condition tab offers 'By field' (a measure, aggregation and comparison) and 'By formula'. By field is the intended no-code path for threshold membership and re-evaluates when data refreshes.

A row-level measure filter would filter marks differently than a dimension condition - know both doors.

Q32 - Correct: B

Objective: 2.6 - Bins

Tableau snaps bins to multiples of the size starting at or below the minimum: bins begin at 20 and run 20-29, 30-39, 40-49, 50-59, 60-69, 70-79 -> six bins cover 22 through 76. Labels display each bin's lower bound.

Counting (76-22)/10 = 5.4 and rounding the wrong way is the intended slip; anchor at the snapped start value instead.

Q33 - Correct: B

Objective: 2.1 - String functions

STARTSWITH (and its twin ENDSWITH) are the dedicated prefix/suffix tests. CONTAINS would wrongly catch a value like "XRET-12", and MID does substring extraction, not testing (though its comparison happens to work, it is not the purpose-built function the question asks for).

Exam phrasing matters: 'begins with' -> STARTSWITH; 'ends with' -> ENDSWITH; 'anywhere' -> CONTAINS.

Q34 - Correct: D

Objective: 2.1 - Date calculations

Three weeks = 21 days; December 20 + 21 days crosses the year boundary cleanly to January 10, 2025. DATEADD handles year/month rollovers natively.

Option B adds three DAYS; option C adds two weeks.

Q35 - Correct: A

Objective: 2.7 - Chart-vs-map vocabulary

The heat map is the crosstab-style grid where Color encodes magnitude across two dimensions. Density and filled maps are geographic; tree maps nest rectangles without a row/column grid.

Common confusion: 'density map' belongs to geography; 'heat map' belongs to tables.

Q36 - Correct: C

Objective: 2.8 - LOD inside an aggregate

The inner LOD stamps every line with its customer's order count (Amy 3, Bo 1, Cy 2); the IF keeps only lines whose customer count exceeds 1; COUNTD then counts those customers: Amy and Cy = 2 (the 'repeat customer' count).

3 is total customers; 6 is total order lines; 1 would count only Amy's tier. This nesting of an LOD inside COUNTD is a high-frequency advanced pattern.

Q37 - Correct: B

Objective: 2.8 - Totals and subtotals

Subtotals are an Analysis > Totals feature that summarizes each outer-dimension block ('Add All Subtotals'). They sit between the detail rows and any grand total, using the measure's default aggregation (configurable).

Reference lines, groups and extra sheets each solve a different, wrong problem here.

Q38 - Correct: A

Objective: 2.8 - Forecast customization

Forecast Options includes 'Ignore last N time periods', built precisely for incomplete trailing periods. Other options you should recognize: forecast length, prediction intervals, aggregation by, and auto/custom trend-and-seasonality models.

Discrete dates or yearly aggregation would break forecasting entirely, not fix the partial month.

Q39 - Correct: A

Objective: 2.6 - Sets in calculations

Sets are booleans: inside a calculation, [Set Name] reads TRUE for members and FALSE otherwise. Wrapping the member's sales in IF...THEN and summing yields the members-only total; IN/OUT can also be shown by dropping the set on Color.

Sets can be used directly in filters, color, set actions, and calculations - the exam checks all four uses.

EXPLANATIONS - DOMAIN 3 - CREATE CONTENT (Q40-Q55, 26%)

Q40 - Correct: C

Objective: 3.1 - Stacked percent chart

Stack the segments, then let the quick table calc re-express each segment as its share of the bar: every region column then totals 100%. This is the standard 'mix' or composition-comparison chart.

The LOD route could force the same result, but the question asks the standard construction - quick table calc on the stacked measure.

Q41 - Correct: B

Objective: 3.2 - Dashboard objects

The Download object adds export buttons (Image / PDF / PowerPoint) directly onto the dashboard canvas - no menu hunting for viewers.

Extensions run custom add-on code; Web Page embeds a URL; Image renders a static picture.

Q42 - Correct: C

Objective: 3.3 - Filter actions

The three behaviors: Show all values (reset to everything), Leave the filtered view (freeze as-is), and Exclude all values (empty the target). Progressive-disclosure designs - 'pick something to see detail' - use Exclude all values.

The three options map exactly onto three different UX patterns; the exam asks each of them.

Q43 - Correct: D

Objective: 3.3 - Highlight actions

All Fields highlights marks matching ANY shared dimension (too broad - the same year/category marks light up too). Selected Fields narrows highlighting to the chosen field(s) only.

Highlighting never removes data; it only emphasizes - that contrast with filter actions is itself exam-favorite material.

Q44 - Correct: A

Objective: 3.3 - Show/hide buttons

Show/hide buttons add an automatic toggle button to a dashboard item (they shine for floating containers holding filters/legends). No calculation is required - that is what separates them from dynamic zone visibility, which needs a boolean field or parameter.

Know the duo: show/hide = manual toggle; DZV = automatic visibility from data/parameters.

Q45 [SELECT TWO] - Correct: B, D

Objective: 3.3 - Parameters on published dashboards

The parameter control is the classic editor; parameter actions let viewers set the parameter by interacting with marks (and both work for web viewers with View rights).

Quick filters drive filters, not parameters; XML editing is an authoring myth in this context.

Q46 - Correct: B

Objective: 3.4 - Default color properties

Default Properties > Color binds member-to-color assignments to the field, so every new view inherits them. Per-sheet edits (option C) are exactly what the question wants to avoid.

The same Default Properties dialog hosts number format, aggregation, sort and comments - know the full family.

Q47 - Correct: D

Objective: 3.1 - Combined/shared axis

Dragging a measure onto another measure's axis merges them into a single shared (blended) axis and Tableau brings in Measure Names on Color. Dual axis keeps two independent scales - the opposite of the requirement.

Shared units + shared scale -> combined axis. Different units/scales -> dual axis (then decide whether to synchronize).

Q48 [SELECT TWO] - Correct: A, C

Objective: 3.3/3.4 - Tooltips

Viz in Tooltip embeds another sheet, filtered to the hovered mark, with a configurable max size; the tooltip editor also accepts any view field plus rich formatting.

Tooltips do not run actions or refresh data - they display context. (The Keep Only/Exclude 'command buttons' inside tooltips can be toggled off, a separate exam line.

Q49 - Correct: D

Objective: 3.1 - Pie chart labels

Duplicating the measure, quick-table-calculating it to Percent of Total, and putting it on Label gives percent labels in four clicks. The LOD route works but is the heavier tool for a display concern.

Remember the pie anatomy too: dimension on Color/Label, measure on Angle.

Q50 - Correct: A

Objective: 3.2 - Dashboard text

Text objects (and worksheet titles/captions) embed live values through the Insert menu - parameters, fields, and workbook metadata. The text re-renders whenever the parameter changes.

This is the standard way to print the current selection in a header; the same Insert trick works in worksheet titles.

Q51 - Correct: C

Objective: 3.2 - Dashboard canvas

The dashboard Grid overlays alignment guides on the canvas and helps snap/snap-align floating and tiled items for pixel-tight layouts.

Containers enforce structure automatically; the grid guides manual placement.

Q52 - Correct: B

Objective: 3.1 - Quick filter dependence

Only Relevant Values makes a quick filter's list respect the other dimension filters in play, so users can never pick an impossible combination. (It costs extra queries - 'Show all values in database' is the fast, context-free alternative.)

This is the dependent/cascading-filter feature every dashboard question about cascading lists targets.

Q53 - Correct: A

Objective: 3.4 - Tooltips off

The Tooltip editor's 'Show Tooltips' checkbox turns them off completely for that worksheet; fields kept on the card then do nothing. Deleting fields leaves an empty popup - the wrong look, not an off switch.

In the same dialog you can also silence the interactive command buttons while keeping passive text.

Q54 - Correct: D

Objective: 3.4 - Mark labels

Tableau deduplicates/culls colliding labels by default; 'Allow labels to overlap other marks' forces them all to draw. The Min/Max, Selected and Highlighted modes are separate selectors for WHICH marks get labels at all.

Overlap-tolerant labeling is the answer to 'some labels randomly missing' complaints.

Q55 - Correct: B

Objective: 3.3 - User guidance

Guiding sentences are plain text: a Text object beside the viz, or the worksheet title/caption. Annotations pin to data, alerts email on thresholds, and quality warnings flag data problems - none instructs users.

Good-analyst instinct the exam rewards: interactivity must be DISCOVERABLE.

EXPLANATIONS - DOMAIN 4 - PUBLISH AND MANAGE CONTENT (Q56-Q60, 9%)

Q56 - Correct: C

Objective: 4.3 - Permissions model

Permissions resolve with a simple rule set: an explicit Deny beats an Allow; nothing allowed means no access. 'Unspecified' is not neutral - it yields no capability, i.e., effectively denied.

Chain to remember: Deny > Allow > Unspecified(=no). Roles cap what permissions can ever grant.

Q57 [SELECT THREE] - Correct: A, B, C

Objective: 4.1 - Exporting content

Image, PDF and PowerPoint are the direct export/print outputs for views and dashboards (crosstab/data exports such as Excel/CSV are the data-side siblings). .hyper and .tfl are data artifacts, never view exports.

On the sheet side also know: Worksheet > Export > Image / Data / Crosstab to Excel, and Copy > Crosstab.

Q58 - Correct: D

Objective: 4.3 - Data-driven alerts

When creating the alert you pick the check cadence. For views on extracts, the evaluation can be tied to data refresh; time-based cadences cap how often emails may fire. Alerts attach to continuous numeric axes only.

Distinguish from subscriptions: alerts are CONDITION-triggered; subscriptions are CALENDAR-triggered snapshots.

Q59 - Correct: A

Objective: 4.3 - Custom views

Custom views store per-user filter/selection state; at save time the author chooses whether the view is visible to everyone or private. Shared custom views then appear in the view switcher for permitted users.

They never change the underlying published content for others - that is their superpower.

Q60 - Correct: A

Objective: 4.2 - Scheduling flows and refreshes

Linked Tasks (Server/Cloud) attach one task to run after another - classically 'run this flow, then refresh that workbook/data source extract' - so the downstream refresh always consumes fresh flow output.

Independent clock schedules race each other and can refresh the workbook off a stale or half-written table.

DOMAIN SCORECARD

Mark your score per domain: Domain 1 __/14 | Domain 2 __/25 | Domain 3 __/16 | Domain 4 __/5. The exam uses compensatory scoring - only the total matters - but a domain below 60% tells you exactly where the next study hour should go. Domain 4 is only 5 questions here, yet those points are the cheapest on the exam; do not give any away.