TABLEAU DATA ANALYST - ADVANCED PRACTICE EXAM 2

ANSWER KEY WITH EXPLANATIONS

QUICK ANSWER KEY

Q1: A	Q11: D	Q21: A	Q31: B	Q41: B, D	Q51: C
Q2: D	Q12: B	Q22: B	Q32: A	Q42: B	Q52: D
Q3: A	Q13: A	Q23: B, D	Q33: A	Q43: A	Q53: B
Q4: B	Q14: A	Q24: B	Q34: C	Q44: C	Q54: C
Q5: D	Q15: D-C-A-F-E-B	Q25: D	Q35: D	Q45: A	Q55: B
Q6: C	Q16: B	Q26: A	Q36: C	Q46: B	Q56: D
Q7: B	Q17: B	Q27: A, D	Q37: B-D-A-C	Q47: B	Q57: B
Q8: B-A-C-D	Q18: D	Q28: D	Q38: D	Q48: B-C-A-D	Q58: C
Q9: A, D	Q19: B	Q29: B	Q39: A, B	Q49: D	Q59: D
Q10: C	Q20: B, C	Q30: C, D, E	Q40: A	Q50: D	Q60: B

Scoring: give yourself 1 point per question only when the answer is COMPLETELY right (all letters on multi-selects; the exact sequence on build lists). 39+/60 keeps you on track for the real exam's 65% pass mark; 45+ means exam-ready.

DETAILED EXPLANATIONS

Question 1 - Correct answer: A

FIXED LODs are computed BEFORE dimension filters but AFTER extract, data source, and context filters. As a plain dimension filter, Category X did not change the FIXED result (East stayed 10 + 20 = 30). The moment the filter is added to context it moves earlier in the order of operations, so the LOD is recomputed over the context-filtered rows only: East = 10. (60 would be double-counting the LOD across two rows; 30 is the pre-context trap.)

Question 2 - Correct answer: D

A blend aggregates the secondary source first and cannot merge at row level; a union stacks rows of identical structure; a relationship keeps tables separate and combines them only at query time, not at row level. A cross-database join adds the Excel connection beside the SQL Server connection on the physical layer and joins them row by row, producing the single denormalized table the question requires.

Question 3 - Correct answer: A

The bullet graph is Show Me's target-comparison chart: a primary-measure bar drawn over shaded qualitative range bands with a reference marker for the comparative (target) measure. A stacked bar and a Gantt bar have no built-in qualitative ranges, and side-by-side circles encode no target.

Question 4 - Correct answer: B

Compute Using Pane (down) partitions the calculation by pane - here each column is a pane, so the percent-of-total denominator is the Consumer column total (100 + 300 = 400), not the table total (1,000). East|Consumer = 100 / 400 = 25%. (10% = 100/1,000 is the Table (down) trap; 16.7% invents a wrong subtotal.)

Question 5 - Correct answer: D

Scheduled refreshes run unattended, so no human is present to answer a credential prompt. The database credentials must be embedded with the connection - set at publish time, or edited afterwards in the data source's Connection settings on Server/Cloud. 'Prompt user' only works for interactive sessions, never for scheduled refreshes.

Question 6 - Correct answer: C

A wildcard union re-scans the folder and pulls in every file whose name matches the pattern - on a live connection those rows arrive in every query. Only stopping the file from matching (move/rename it or refine the pattern) guarantees exclusion. An extract filter would require an extract, which is disallowed; Data Interpreter cleans formatting, not file inclusion.

Question 7 - Correct answer: B

Unlike FIXED, INCLUDE and EXCLUDE expressions are computed at the viz level AFTER dimension filters apply. Hiding S3 reduces the visible East rows to 100 + 200 = 300, and the EXCLUDE region total recomputes over those rows. (600 is the FIXED-behavior trap - only FIXED sits earlier in the order of operations.)

Question 8 - Correct answer: B-A-C-D

A waterfall is a Gantt chart: the running total positions each mark vertically, and the negated measure sizes the Gantt bar backward from that position so the visible bar spans only the quarter's change. Build the running total first (B), switch to Gantt (A), size with the negated measure (C), then verify/format (D).

Question 9 - Correct answer: A, D

In Specific Dimensions, the checked fields define the addressing (the walk direction); unchecked fields define the partitioning (where the calc restarts). Ticking every dimension leaves nothing to partition on, so the calc runs across the entire table in one partition. Table calculations never rewrite source rows, and Specific Dimensions exists precisely for complex multi-dimension views - it is never disabled by having dimensions on both shelves.

Question 10 - Correct answer: C

An Aggregate step groups rows that share identical values in the grouped fields. Grouping by every field makes each unique row-signature its own group, so exact duplicates collapse into one. Trimming spaces fixes formatting variance, not identical duplicates; a self-join multiplies rows; pivoting reshapes columns and never deduplicates.

Question 11 - Correct answer: D

A custom number format uses the pattern POSITIVE;NEGATIVE (optionally ;ZERO). Setting #,##0;(#,##0) as the field's default number format makes negative parentheses appear everywhere with no calculated field and no data change. A string calc destroys the numeric type, and Tableau offers neither an axis 'Accounting' style nor a parenthesized currency data type.

Question 12 - Correct answer: B

The FIXED LOD stamps each East city with 300 and each West city with 600. Averaging across the six marks: (300 + 300 + 600 + 600 + 600 + 600) / 6 = 3,000 / 6 = 500. The trap is 450 - the simple average of the two region totals - but the view has six marks, so West's four cities weight its region total four times as heavily.

Question 13 - Correct answer: A

File -> Export As Version downgrades a workbook to an older format so it can be opened by - and published to - an older Server. There is no 'Older Server format' choice in Save As; switching .twbx to .twb keeps the same internal version; and Tableau Cloud is not a bridge to an older on-prem Server.

Question 14 - Correct answer: A

A blend aggregates the secondary source to the level of the linking fields in the view. The view reaches only State grain, but each State maps to multiple SalesRep values, so the field cannot resolve to a single value and Tableau displays *. The fix is to bring SalesRep itself into the view (finer granularity) or rework the linking relationship.

Question 15 - Correct answer: D-C-A-F-E-B

The pipeline is: extract filters -> data source filters -> context filters -> FIXED LODs -> dimension filters -> INCLUDE/EXCLUDE LODs (viz level) -> measure filters -> table calculations -> table calculation filters. This is why a FIXED LOD ignores a dimension filter (C before A) but an EXCLUDE LOD does not (F after A), and why a table-calculation filter (B) leaves every computed number underneath untouched.

Question 16 - Correct answer: B

Stepped Color chops a continuous measure into N discrete color blocks - exactly three flat bands here. Sequential and diverging palettes interpolate continuously (a gradient), and a diverging center only pins where the midpoint sits. Sorting changes mark order, never the color encoding.

Question 17 - Correct answer: B

Measure filters are applied BEFORE table calculations run. 'At least 500' removes South and East from the view's data, and the percent-of-total then computes over the single remaining mark: 500 / 500 = 100%. This is the mirror-image trap of a table-CALCULATION filter, which hides marks without recomputing totals (then North would stay 50%).

Question 18 - Correct answer: D

Default Properties -> Fiscal Year Start sets the fiscal calendar for a date field once; every view, date hierarchy, and date part (quarter, year) then uses fiscal labeling automatically. A shifting calculated field moves real dates and corrupts calendar analysis, axis formatting cannot relabel fiscal quarters, and importing derived columns is unnecessary maintenance.

Question 19 - Correct answer: B

POSITIVE offsets look FORWARD along the addressing direction: WINDOW_AVG(..., 1, 2) averages the next two marks after the current one. February averages March (25) and April (35) = 30. (15 is February itself - the offsets exclude the current row; 20 averages looking backward - the -2,-1 trap; NULL occurs only when fewer than two forward marks exist, e.g., on May.)

Question 20 - Correct answer: B, C

Pages add a slider with play controls that steps the view through each member (e.g., month by month), and the Show history option keeps prior marks on screen with fading trails - the classic animation setup. Pages work on ANY view, maps included but not required, and they filter the display only: they never modify the underlying data.

Question 21 - Correct answer: A

FIND(string, substring, start) searches from the 'start' position onward. Positions: A=1, B=2, first hyphen=3, C=4, D=5, second hyphen=6. Starting at 4 skips the first hyphen and finds the second at position 6. (3 is found only without the start argument; 0 would mean 'not found'.)

Question 22 - Correct answer: B

Tableau evaluates permissions in a fixed order: user rules beat group rules; among group rules, ANY explicit Deny beats any Allow; if no rule matches, the fall-back default applies. Here two group rules conflict, so the Deny wins. (The site role still caps everything regardless - a Viewer could never save.)

Question 23 - Correct answer: B, D

ISDATE validates REAL calendar dates: February 30 never exists, so FALSE - that statement is true. 2024 is a leap year, so February 29, 2024 is valid and MAKEDATE returns it - also true. For the same leap-year reason, one day before March 1, 2024 is February 29, NOT February 28, so the DATEADD statement fails. And DATENAME returns a STRING ('October'); it is DATEPART that returns the integer 10.

Question 24 - Correct answer: B

Label options include labeling 'Selected' marks - labels appear when a mark is clicked and disappear when it is deselected. 'All' labels every mark (the density problem), Min/Max labels only extremes, and tooltips are hover text, not persistent on-chart labels.

Question 25 - Correct answer: D

Clicking the '14 unknown' indicator (or Map -> Edit Locations) opens the Edit Locations dialog: you can set a fixed Country/Region to remove ambiguity and manually match misspelled values to known locations. Geographic role changes do not fix name matching, Data Interpreter cleans table structure, and grouping hides rather than resolves the problem.

Question 26 - Correct answer: A

Tableau's ROUND rounds halves AWAY FROM ZERO: ROUND(2.5, 0) = 3 and ROUND(-2.5, 0) = -3. This pairs with the classic INT() trap (INT(-2.5) = -2, truncating toward zero) - examiners love testing that rounding and truncation travel in opposite directions for negative numbers.

Question 27 - Correct answer: A, D

A published data source arrives as a finished model: its physical-layer tools are locked, so unions and pivots (which reshape the physical schema) are unavailable - you would rebuild them in Prep or against the original connection. Data Interpreter is Excel/text-only, so its claimed universal availability is the trap, and the published relationship model likewise cannot be edited from Desktop.

Question 28 - Correct answer: D

IFNULL([Ship Date], TODAY()) substitutes today's date when Ship Date is missing, and DATEDIFF('day', start, end) then measures elapsed days for both shipped and in-transit orders. ZN converts nulls to ZERO - meaningless on dates; plain subtraction leaves NULLs that pollute aggregates; and nothing is ever EQUAL to NULL, so an '=' test can never detect the missing dates.

Question 29 - Correct answer: B

The Density mark type overlays a kernel-density heat layer on ANY quantitative view, not just maps. Trend lines add fit lines but don't declutter, and higher opacity worsens overplotting. A filled map requires geographic fields this view doesn't have.

Question 30 - Correct answer: C, D, E

MAX over a FIXED region total returns the biggest regional sum: the FIXED result is stamped per row and the outer MAX picks the largest of them. Nested LODs evaluate inside-out - the inner expression computes first and the outer one aggregates its results. A row-level FIXED result such as a customer's first-ever order date can absolutely serve as a dimension. The two decoys: INCLUDE results are row-level values normally consumed as measures - never forced to be dimensions; and EXCLUDE is computed at viz level AFTER context filters have been applied, so it never ignores them.

Question 31 - Correct answer: B

Interacting requires only the 'View' capability. 'Web Edit' controls browser editing, and 'Download Workbook/Save a Copy' controls taking the .twbx file - denying those two preserves viewing while blocking edit and export. Denying 'View' hides the content entirely, a story does not change capability enforcement, and republishing changes nothing.

Question 32 - Correct answer: A

Relative to First fixes the comparison base at the partition's first mark (Jan = 80) for every mark: March = (120 - 80) / 80 = 50%. (20% is the 'Relative to Previous' trap - (120-100)/100; 25% is February's own value; NULL applies only to January, the reference mark itself.)

Question 33 - Correct answer: A

A bump chart plots rank over time: Rows hold RANK(SUM([Sales])) computed so that within each month partition the products are ranked against one another (Specific Dimensions: Product), with Product on Color/Detail so each product becomes a line. Raw SUM never shows position swaps, running totals accumulate rather than rank, and INDEX() merely numbers rows in view order - not by sales rank.

Question 34 - Correct answer: C

With hard joins, Tableau normally queries both tables to honor join semantics (unmatched or duplicated rows). Telling Tableau to Assume Referential Integrity guarantees every Orders row matches exactly one Customers row, so when a view touches only Orders fields the optimizer can drop the right table entirely (join culling). Hiding fields doesn't stop the queries; unions and second connections change the model, not the culling.

Question 35 - Correct answer: D

'First-ever order' lives at customer grain, so you need a customer-level LOD: {FIXED [Customer] : MIN([Order Date])} per customer, wrapped in YEAR() and used as a dimension (the cohort year), then count DISTINCT customers per cohort. YEAR([Order Date]) bins orders by their own year (repeat buyers look 'new'), a running total accumulates instead of counting per year, and plain MIN is view-level, not customer-level.

Question 36 - Correct answer: C

Range size lets the dashboard scale with its container while clamping it between the specified minimum and maximum - exactly the requirement. Automatic has no bounds at all; Fixed never grows; Fit to Width is a sheet-level scaling option, not a dashboard size with min/max bounds.

Question 37 - Correct answer: B-D-A-C

The parameter must exist before anything can reference it (B). The CASE calculation then turns the selected grain into a real date value - each branch truncates appropriately (D). Only then can the field go on the axis with the parameter control exposed (A), and the final step verifies behavior (C). Placing a field that doesn't exist yet breaks the chain.

Question 38 - Correct answer: D

The Salesforce connector queries through Salesforce's APIs (objects with their own join model), not raw SQL, so the Custom SQL entry point is unavailable for that connection type. Data source filters or a Tableau Prep step are the supported ways to subset rows. Editions, operating systems, and extracts have nothing to do with it.

Question 39 - Correct answer: A, B

Sets are usable as row-level booleans - IN or OUT - so IF [Set] THEN ... END is valid, and set actions exist precisely to mutate membership from clicks, cascading into other sheets. The decoys: a set belongs to exactly ONE dimension, and constant sets hold a frozen member list - only condition/top-N dynamic sets recompute on refresh.

Question 40 - Correct answer: A

Every dashboard item exposes Outer Padding (plus border/background) in the Layout pane - 8 px outer padding on the items produces uniform gutters declaratively. Blank objects add unmanaged tiles that break alignment on resize; the Desktop-side chart area doesn't control dashboard-level gaps; exporting to PowerPoint abandons the interactive dashboard entirely.

Question 41 - Correct answer: B, D

Bins materialize as discrete dimensions (the blue pill) that bucket a continuous measure, and they behave like other dimensions in views. Tableau's suggested bin size comes from the field's range and record distribution. The decoys: bins can only be created on measures (never strings), and resizing a bin layers a grouping ON TOP of the data without ever editing it.

Question 42 - Correct answer: B

Hidden fields still exist - they are merely removed from display. The Data pane menu's 'Show Hidden Fields' reveals them greyed out so individual fields can be unhidden. Refreshing doesn't unhide anything, and nothing was deleted from the connection - reconnecting from scratch is an unnecessary rebuild.

Question 43 - Correct answer: A

The INCLUDE expression evaluates at Region + Customer level first: customer totals A = 150 and B = 200. Wrapping it in COUNT then aggregates those per-customer values inside the East mark: COUNT(150, 200) = 2 - effectively the number of customers. (350 is SUM; 3 counts the underlying order rows; 150 is AVERAGE.)

Question 44 - Correct answer: C

URL actions can target either a new browser tab or a Web Page object on the same dashboard; choosing the object as the URL Target renders the field-templated URL (with inserted) in place. Navigation buttons jump between dashboards/sheets, not URLs, and filter actions never touch web content.

Question 45 - Correct answer: A

Tableau never sums the displayed marks for a grand total - it recomputes the aggregation over the whole underlying partition. COUNTD across all rows is 70 because 35 customers were counted once in East and again in West or South. This non-additivity of distinct counts is a classic exam trap: the total is correct, not broken.

Question 46 - Correct answer: B

Incremental refresh compares the chosen column against the maximum already in the extract and pulls only rows beyond it. Updated historical rows keep their original dates, so they are never re-fetched. A full refresh reloads everything and picks up corrections (a common periodic practice for exactly this reason). The refresh IS succeeding - its semantics, not the network, exclude changed old rows.

Question 47 - Correct answer: B

Worksheet -> Copy -> Formatting captures every format setting; Paste > Formatting applies it to the other sheets in seconds (tabs can be multi-selected with Ctrl-click). Manual rebuilding and hand-editing XML are slow and error-prone, and duplicating the sheet 12 times means rebuilding the actual views anyway.

Question 48 - Correct answer: B-C-A-D

Define the parameter first (B), then the boolean test comparing each product's rank against it (C). Because RANK is a table calculation, the filter pill needs Compute Using set so products are ranked in the intended direction/partition before being kept on Filters as True (A); only then validate with the control (D). Note this uses a table-calculation filter - the underlying data stays intact and other marks/LOD math are unaffected.

Question 49 - Correct answer: D

Incrementality is a property of the extract itself: the column identifying new rows (an increasing date or ID) must be defined in the extract's configuration in Desktop, and the data source republished. A schedule merely triggers whatever refresh modes the published extract supports - you cannot bolt incrementality onto an extract that has no incremental column defined.

Question 50 - Correct answer: D

Range is designed for exactly this: a bounded continuum with a step size that renders as a slider or type-in box enforcing the bounds and increments. A manually entered List is finite and unmaintainable; 'All' accepts unbounded junk with no min/max/step control; and field-driven lists are for categorical pickers, not numeric sliders.

Question 51 - Correct answer: C

Prep's wildcard pivot scans the input at run time and pivots every column whose name matches the search pattern - new month columns flow through untouched. Manual field selection freezes the column list at design time; a join combines on keys and never unpivots; and a union stacks tables, not columns.

Question 52 - Correct answer: D

When a dimension is in the view, each mark corresponds to exactly one Region value, but Tableau still requires dimension references to be aggregated inside an aggregate expression. ATTR([Region]) returns the mark's single Region value (or '*' if several existed), satisfying the aggregate/non-aggregate rule. Dropping SUM would silently switch to row-level logic, the naive brace syntax is not a valid LOD around an IF, and ISNULL has nothing to do with aggregation.

Question 53 - Correct answer: B

'Distribute Contents Evenly' on a layout container forces its children to share the container's space equally and keeps them equal as the dashboard resizes. Manual dragging is approximate and doesn't survive resizing; Blank objects waste space without enforcing equality; floating locks positions but loses the automatic, proportional resizing that makes tiled layouts attractive.

Question 54 - Correct answer: C

Hide Sheet (the grey italic tab) keeps a worksheet fully functional inside dashboards, stories, and actions while excluding it from the published tab list. Deleting the sheets destroys the dashboards they feed; helper sheets cannot live in another workbook while powering this one; and a data source filter changes data, not sheet visibility.

Question 55 - Correct answer: B

Automatic blending links only fields with identical names. Edit Blend Relationships (or the orange link icon) lets you define custom field pairings across the two sources - here Region to Territory - after which the secondary aggregates correctly at Region level. Renaming a file changes nothing about field names, unions stack rows for a different purpose, and placing both fields side by side doesn't teach Tableau they represent the same entity.

Question 56 - Correct answer: D

DATEDIFF('day', start, end) counts day boundaries from start to end - conceptually end minus start. Here the START (March 10) is later than the END (March 1), so the result is negative: 1 - 10 = -9. Reversed arguments flip the sign; they never raise an error - one of the exam's favorite date traps.

Question 57 - Correct answer: B

'Show Search Box' lets users type to locate members inside a long list, and 'Show Apply Button' defers the query until they finish ticking and click Apply - one round-trip instead of one per tick. Wildcard-only hides the list entirely, single-value forbids multi-picking, and context changes WHEN the filter applies in the pipeline, not how often the control fires queries.

Question 58 - Correct answer: C

Unions generate metadata columns - Table Name (and File Path for file unions) - identifying each row's source file. A calculated field such as SPLIT([Table Name], '_', 2) extracts 'East' as a real Region dimension. No manual file edits are needed, merging mismatched fields only aligns columns, and a join cannot fabricate per-file metadata.

Question 59 - Correct answer: D

Parameters act as constants, so comparing the aggregated SUM([Profit]) with [Profit Target] is a legal, fully aggregate comparison whose result colors each mark - expose the parameter control and the colors update live. The braced variant is invalid LOD syntax (no keyword/dimensions); the inverted version both flips the logic and, lacking an ELSE, returns NULL for misses; and Tableau has no COLOR() function.

Question 60 - Correct answer: B

Percent Difference From previous = (current - previous) / ABS(previous). LOOKUP(..., -1) reaches the previous mark, and the denominator must be ABS(previous) so results stay directionally correct even when the base is negative. Dividing by SUM([Sales]) uses the current value as the denominator; the +1 variant looks FORWARD instead of back; and the RUNNING_SUM variant divides an accumulating total by the current value - none equal the quick table calculation.