SALESFORCE CERTIFIED TABLEAU DATA ANALYST (ANALYTICS-DA-201)

ADVANCED PRACTICE EXAM - 60 QUESTIONS

Questions	60 (mirrors the real exam's 60 scored items)
Time allowed	105 minutes - same as the real exam
Passing score	65% = 39 / 60 correct (this set is harder - aim for 45+)
Question types	Single choice, [SELECT TWO / THREE] multi-select, and [BUILD LIST] sequencing
Scoring	No partial credit - a multi-select or build-list item counts only when every part is exact
Navigation	Flag and return, as in the real exam; there is no penalty for guessing - answer everything

How to use this paper: this is a deliberately HARDER tier than the real exam, built to stretch you past the pass mark. It is heavily loaded with the areas you flagged as weak: Level-of-Detail expressions (including nested and filter-interplay traces), calculation traces you must evaluate by hand, table-calculation edge cases, multi-select items, and step-sequencing (build list) questions. Trace every calculation on paper - do not guess.
Conventions: unless marked otherwise each question has exactly ONE correct answer. Items marked [SELECT TWO] / [SELECT THREE] need EXACTLY that many correct choices - no partial credit. [BUILD LIST] items are correct only when the full letter sequence is right (write it like B-D-C-A). All questions are new - none are repeated from your earlier papers.

ANSWER EVERYTHING FIRST - THE ANSWER SHEET IS AT THE END. THE ANSWER KEY (advanced-answers.pdf) IS IN THE ANSWERS FOLDER - DO NOT OPEN IT UNTIL YOU FINISH.

Question 1. A security policy requires that, at the moment an analyst connects to a Snowflake data warehouse, a temporary session table named #UserContext is created once - it stores the analyst's user ID for row-level security - and every query Tableau then runs must be able to reference it. The underlying Orders table itself must be queried exactly as it is stored. Which connection feature should you use?

A. Write the connection as Custom SQL that begins with CREATE TABLE #UserContext, then SELECT * FROM Orders

B. Use Initial SQL on the connection so #UserContext is created once, when the connection opens

C. Apply a data source filter that references the current user

D. Create an extract with a filter on User ID

Question 2. An analyst creates an extract of a five-year orders table. In the Extract dialog she leaves 'Aggregate data for visible dimensions' selected and rolls the date field up to Years. Three weeks later, marketing asks for a monthly seasonality view built on this extract. What happens when she drags a Month-level date into a new worksheet?

A. Monthly values appear - extracts always keep full row-level detail

B. The month values all show as January, but the data is correct

C. Month-level detail does not exist in the extract; she must re-create the extract without rolling dates up (or with Month visible) to get that grain

D. She can restore monthly detail by running an incremental refresh

Question 3. Customers (left table) is joined to a Marketing_Preferences table on Email. 500 customer rows have NULL in Email, and Marketing_Preferences contains 120 rows whose Email is also NULL. After the left join, how do the NULL-email customers appear - and can they ever match the NULL-email preference rows?

A. All 500 appear with NULL preference fields; NULL join keys never match other NULLs, so no pairing is possible

B. All 500 match the 120 NULL rows, creating duplicated rows

C. The 500 rows are dropped from the join result entirely

D. NULL matches NULL in Tableau joins, so all 500 customers appear with preferences attached

Question 4. You union two regional CSV exports. Region A's file names its key column CustID; Region B's file uses CustomerID for the same logical field. Everything else is identical. What does the unioned table look like?

A. Tableau automatically detects that CustID and CustomerID are the same field and merges them

B. The result contains CustID and CustomerID as two separate columns, each NULL on the other file's rows; rename one field so they align

C. The union fails with a schema-mismatch error

D. The first file's schema wins and Region B's key column is dropped

Question 5. Primary source: SUM(Sales) by Region. Secondary source: a published Targets data source with one row per Region per Month (East: 10,000 for each of Jan, Feb, Mar). The sources are blended and linked on Region; Month is NOT in the view. SUM([Target]) is added to the East mark. What does it display?

A. 10,000

B. 30,000

C. NULL - a blend cannot aggregate the secondary source

D. An error - a blend requires Month to be in the view

Question 6. A workbook uses a live connection to SQL Server. Users pick a Region in a parameter, and you want the parameter's value to flow into the SQL query itself, so the database returns only that region's rows each time the parameter changes. Which feature supports this?

A. A data source filter that references the parameter

B. A quick filter set to 'Only relevant values'

C. Writing the SELECT statement as Custom SQL that embeds the parameter

D. A context filter driven by the parameter

Question 7. A colleague needs the governed Sales data source AND its data as one portable file so she can work fully offline on a plane, with no access to Tableau Server. Which file type should you send her?

A. .twb - the workbook file with the data embedded

B. .tds - the data source file

C. .tfl - the Tableau Prep flow file

D. .tdsx - the packaged data source, which bundles the connection metadata with its .hyper extract

Question 8. A Tableau Prep flow's input step loads a 5-million-row table. The profile pane shows a 10,000-row sample (Quick select). How many rows does the flow write to its output when the flow runs?

A. Approximately 10,000 - the sample is the flow's working set

B. All 5 million - sampling only affects the profile/preview panes; the full data is processed at run

C. A random 1% (50,000 rows)

D. The flow fails at run time until sampling is turned off

Question 9. [SELECT TWO] Which TWO practices genuinely improve data source and workbook performance?

A. Adding every dimension as a multiple-value quick filter on every sheet

B. Hiding unused fields in the data source

C. Turning off aggregation so views plot every underlying row

D. Replacing a slow live connection with an extract for large sources

Question 10. [BUILD LIST] You receive a wide Excel report: a title banner, merged cells, subtotal rows, footnotes, and twelve monthly revenue columns (Jan ... Dec). You need a tidy Month/Revenue table to chart the trend. Arrange the steps in the correct order.

A. Rename the generated pivot field names (e.g., Month, Revenue) and build the line chart

B. Connect to the Excel file in Tableau Desktop

C. Select the twelve month columns on the data source page and pivot them into rows

D. Turn on the Data Interpreter so titles, subtotals, and footnotes are removed and headers detected

Write the correct sequence of letters, e.g., B-D-C-A.

Question 11. A Budget table stores one row per Region per [Budget Year] (an integer such as 2024). The Orders table stores [Order Date] (a true date). You must relate Budget to Orders so each region-year budget pairs with that year's orders. How should you set this up?

A. Relate the tables directly: [Order Date] = [Budget Year]

B. Create a calculated field STR([Budget Year]) in Budget and relate it to [Order Date]

C. Relate the tables on [Region] only and ignore the year mismatch

D. Create a calculated field YEAR([Order Date]) in the Orders logical table, then define the relationship between that field and [Budget Year]

Question 12. A live Oracle dashboard fires hundreds of queries per session and the DBA reports heavy repeated table scans. Every sheet always filters to the current fiscal year first, then applies many more filters on top. Which single change reduces the database's repeated scanning work?

A. Add more dimension filters to each sheet

B. Convert the measure filters into table calculation filters

C. Promote the Fiscal Year filter to a context filter, so a temporary filtered subset is built once and reused by all subsequent queries

D. Sort every sheet descending by Sales

Question 13. [SELECT THREE] Which THREE statements about Tableau's logical (relationship) data layer are TRUE?

A. Related tables must contain identically named key fields - relationship clauses cannot be edited or use calculated fields

B. Double-clicking a logical table on the data source page opens its physical join canvas, where additional tables can be joined at that table's grain

C. When a viz uses measures from two related logical tables, each table's measure is computed at its own level and the results are combined at the viz's grain - so SUM is not inflated by the other table's rows

D. A relationship clause can be defined on calculated fields and on multiple pairs of fields

Question 14. A table of 3,000 stores contains numeric [Latitude]/[Longitude] columns. A second table contains sales-territory polygons. Each store must be tagged with the name of the territory that contains it. Which approach should you use?

A. Join the two tables on territory name

B. Union the two tables and blend on Latitude

C. Create a physical join between the tables using a spatial Intersects clause (building points from lat/long)

D. Create a combined set of stores and territories

Question 15. A view shows Region and Customer on rows. You add the calculation {EXCLUDE [Region] : SUM([Sales])}. What value appears on the West | Cara mark?

Region	Customer	Sales
East	Anna	50
East	Ben	70
East	Cara	60
West	Cara	90
West	Dana	20

A. 90

B. 290

C. 60

D. 150

Question 16. Using the same data as the previous question: a view shows Customer only. Measure M = {FIXED [Region], [Customer] : SUM([Sales])}, displayed with AVG. A normal (non-context) dimension filter keeps Region = East only. What does the Cara mark show for AVG(M)?

A. 75

B. 90

C. 60

D. 150

Question 17. A retention analysis needs each order's customer tenure: the number of year boundaries between the customer's FIRST-ever order and that order. One order was placed 2024-06-15 by a customer whose first order was 2021-09-01. Which calculation is correct, and what does it return?

A. DATEDIFF('year', [Order Date], {FIXED [Customer] : MIN([Order Date])}) - it returns 3

B. DATEDIFF('year', {FIXED [Customer] : MAX([Order Date])}, [Order Date]) - it returns 3

C. DATEDIFF('year', {FIXED [Customer] : MIN([Order Date])}, [Order Date]) - it returns 3

D. DATEDIFF('year', YEAR(MIN([Order Date])), YEAR([Order Date])) - it returns 3

Question 18. A view shows one mark per Region. The field {FIXED [Region] : AVG({FIXED [Region], [Customer] : SUM([Sales])})} is added. Using this data, what does the West mark display?

Region	Customer	Sales
East	A	100
East	B	200
West	C	50
West	D	100
West	E	150

A. 300

B. 150

C. 50

D. 100

Question 19. A bar chart of 10 states is sorted descending by SUM(Sales). INDEX() and FIRST() are both computed Table (Down). What values appear on the 5th state's mark?

A. INDEX() = 5 and FIRST() = 4

B. INDEX() = 4 and FIRST() = -5

C. INDEX() = 5 and FIRST() = -4

D. INDEX() = 5 and FIRST() = -5

Question 20. Monthly SUM(Sales): Jan 10, Feb 20, Mar 30, Apr 40. The calculation WINDOW_AVG(SUM([Sales]), -1, 1) computes Table (across). What does March show?

A. 20

B. 40

C. 90

D. 30

Question 21. A view shows monthly sales: Jan = 100, Feb = 50, Mar = 200, computed Table (across). The calculation LOOKUP(SUM([Sales]), FIRST()) is added. What does the March mark display?

A. 200

B. 100

C. 50

D. 350

Question 22. Four products have SUM(Sales) of 500, 400, 400, and 300, ranked descending. Management wants the ranking sequence 1, 3, 3, 4 - tied items receive the WORST (highest) of the tied ranks. Which function produces exactly that?

A. RANK(SUM([Sales]), 'desc')

B. RANK_DENSE(SUM([Sales]), 'desc')

C. RANK_UNIQUE(SUM([Sales]), 'desc')

D. RANK_MODIFIED(SUM([Sales]), 'desc')

Question 23. What does DATEADD('month', 1, #2024-01-31#) return?

A. #2024-02-28#

B. #2024-02-29#

C. #2024-03-02#

D. An error, because February has fewer than 31 days

Question 24. What does DATEDIFF('month', #2024-01-31#, #2024-02-01#) return?

A. 0

B. NULL, because a full month has not elapsed

C. 1

D. Approximately 0.03 - the elapsed fraction of a month

Question 25. Order codes look like US-2024-1186. Which expression returns the LAST token (1186) regardless of how many tokens precede it?

A. SPLIT([Code], "-", -1)

B. RIGHT([Code], 4)

C. SPLIT([Code], "-", 3)

D. MID([Code], FIND([Code], "-"), 4)

Question 26. What does ROUND(4567.89, -2) return?

A. 4500

B. 4567.89

C. 4600

D. 4700

Question 27. Which expression returns -4 for the input -3.2?

A. CEILING(-3.2)

B. INT(-3.2)

C. ROUND(-3.2, 0)

D. FLOOR(-3.2)

Question 28. Some rows have NULL in [Ship Date]. What does this calculation return on those rows? IF [Ship Date] > [Order Date] THEN "Late" ELSE "On time" END

A. NULL

B. "Late"

C. An error - the calculation stops

D. "On time"

Question 29. A daily view (one mark per day) must show each day's SUM(Sales) against the AVERAGE of all the daily totals in the view (not the average of the underlying rows). Which calculation returns that average daily total?

A. AVG([Sales])

B. WINDOW_AVG(SUM([Sales]), FIRST(), LAST())

C. {FIXED : AVG([Sales])}

D. RUNNING_AVG(SUM([Sales]))

Question 30. [SELECT TWO] Which TWO statements about parameters are TRUE?

A. A parameter can hold several selected values at the same time

B. When a parameter's list comes from a field, it can be set to refresh each time the workbook is opened, so new values appear without editing the parameter

C. When a parameter action runs on a selection of several marks, it can assign an aggregated value (such as MAX of the clicked field)

D. Changing a parameter's value triggers an extract refresh before the view recalculates

Question 31. A view shows AVG([Sales]) by Region and row grand totals are switched on. What does the default grand total display?

A. The average of the displayed regional averages

B. The sum of the regional averages

C. NULL - averaged measures cannot be totaled

D. The average computed across all underlying rows, as if AVG were run over the entire data set

Question 32. Which expression returns the distance between two stored latitude/longitude pairs in kilometers?

A. DISTANCE([lat1], [lon1], [lat2], [lon2], 'km')

B. DISTANCE(MAKEPOINT([lat1], [lon1]), MAKEPOINT([lat2], [lon2]), 'km')

C. DISTANCE(MAKEPOINT([lat1], [lon1]), MAKEPOINT([lat2], [lon2]))

D. DISTANCE(MAKELINE([lat1], [lon1]), MAKELINE([lat2], [lon2]), 'km')

Question 33. Using Tableau's default date properties (weeks start on Sunday), what does DATETRUNC('week', #2024-10-16#) return? (Oct 16, 2024 is a Wednesday.)

A. #2024-10-16# - the date itself

B. #2024-10-13# - the Sunday that begins that week

C. #2024-10-14# - the Monday that begins that week

D. #2024-10-20# - the Sunday that ends that week

Question 34. January orders (each row is one order) are shown below. The view displays one mark per month with MAX({INCLUDE [Order Date] : COUNTD([Customer])}). What does January display?

Order Date	Customer
Jan 1	Anna
Jan 1	Ben
Jan 2	Ben
Jan 2	Cara
Jan 3	Cara

A. 1

B. 5

C. 2

D. 3

Question 35. [SELECT TWO] Which TWO statements about aggregate functions are TRUE?

A. COUNT([Email]) counts only the rows where Email is not NULL

B. COUNTD() treats NULL as a distinct member and counts it

C. SUM() over a group in which every row is NULL returns NULL, not zero

D. MEDIAN() behaves identically on every live database connection

Question 36. A view shows a single mark for the customer Anna, whose orders span two regions (East and West); Region is not in the view. What does ATTR([Region]) display on that mark?

A. "East"

B. "East, West"

C. "*"

D. NULL

Question 37. [SELECT THREE] Which THREE statements about date functions are TRUE?

A. DATETRUNC('month', #2024-04-19#) returns #2024-04-01#

B. DATENAME('weekday', #2024-10-16#) returns an integer such as 3

C. NOW() returns the current date and time; TODAY() returns the current date with no time component

D. DATEADD('day', -7, #2024-06-10#) returns #2024-06-03#

Question 38. [BUILD LIST] Tableau applies filters and computes calculations in a strict order. Arrange these stages from FIRST applied to LAST applied.

A. Context filters are applied

B. Table calculation filters are applied

C. Extract filters are applied

D. LOD expressions (FIXED / INCLUDE / EXCLUDE) are computed

E. Dimension filters are applied

F. Data source filters are applied

G. Measure filters are applied

Write the correct sequence of letters, e.g., C-F-A-D-E-G-B.

Question 39. [BUILD LIST] Marketing wants customers grouped by the year of their FIRST order (an acquisition cohort). Arrange the steps in the correct order.

A. Drag Acquisition Year onto the view and add SUM([Sales]) to compare cohorts

B. Create the calculated field {FIXED [Customer] : MIN([Order Date])}

C. Wrap that LOD in YEAR(...) and save the result as Acquisition Year

D. Spot-check that each customer's earliest order lands in the correct cohort year

Write the correct sequence of letters, e.g., B-C-A-D.

Question 40. [BUILD LIST] Configure a dashboard action so that HOVERING over a bar in the overview sheet filters the detail sheet, on the Region field only (no other shared fields). Arrange the steps in the correct order.

A. Set Run action on: Hover

B. Build the overview and detail worksheets and place them on the dashboard

C. In the action dialog, choose Source Sheets = overview and Target Sheets = detail

D. Dashboard menu -> Actions -> Add Action -> Filter

E. Under Target Filters, choose Selected Fields and add only Region

F. Hover over a bar in the overview and confirm the detail sheet responds

Write the correct sequence of letters, e.g., B-D-C-A-E-F.

Question 41. A dashboard must scale its content so the whole dashboard fills the browser canvas in BOTH directions - no scrolling and no letterboxing. Which setting fits?

A. Size: Fixed, with Fit: Width

B. Size: Automatic, with Fit: Entire View

C. Size: Automatic, with Fit: Normal

D. Size: Range, with Fit: Height

Question 42. On a customer dashboard, hovering over a bar must show a mini monthly trend for THAT customer inside the tooltip. Which configuration makes the embedded tooltip sheet respond to the hovered bar?

A. Insert the trend sheet into the tooltip and hide its x-axis - filtering happens automatically

B. Add a highlight action from the bar chart to the trend sheet

C. In the tooltip editor, insert the trend sheet and enable the tooltip filter option (Filter: All fields, or the targeted Customer field)

D. Place the trend sheet on the dashboard, then drag the sheet object into the tooltip

Question 43. [SELECT TWO] Which TWO statements about floating dashboard items are TRUE?

A. Floating items can overlap other floating items

B. Floating items cannot contain worksheets - only shapes and text

C. A dashboard must be either all tiled or all floating; the two cannot be mixed

D. A floating container can hold worksheets and objects such as text, images, and filter cards

Question 44. On a scatter plot you must attach a short text callout, with an arrow, to ONE specific dot, and the callout must move if the mark moves. Which annotation type should you use?

A. Area annotation

B. Point annotation

C. Title annotation

D. Mark annotation

Question 45. A worksheet title must always display the currently selected value of the Region Selector parameter, updating live as users change it. How do you build this?

A. Type the parameter's current value into the title text - it will update automatically

B. Edit Title -> position the cursor -> Insert -> Parameters -> Region Selector

C. Add a text object above the sheet on a dashboard and link it to the parameter

D. Rename the worksheet tab to match the parameter

Question 46. A Profit Ratio measure is colored with a red-blue palette. Negative ratios must be exactly red and positives exactly blue, with the boundary pinned at zero, so near-zero positives do not wash toward red. How do you configure the color?

A. Edit Colors -> choose Sequential -> enable Use Full Color Range

B. Edit Colors -> choose Diverging -> Advanced -> set the Center to 0

C. Sort the view descending so positives render first

D. Edit Colors -> choose Sequential -> enable Stepped Color

Question 47. An executive complains that the profit axis shows unreadable numbers like 1,500,000. You want them displayed as 1.5M without changing the data. What do you do?

A. Divide Profit by 1,000,000 in a calculated field and re-aggregate

B. Change the Profit field's data type to string and append "M"

C. Right-click the axis -> Format -> Scale (Numbers) -> set the display units to Millions

D. Set the axis to a logarithmic scale

Question 48. [SELECT THREE] Which THREE objects can be added directly to a dashboard?

A. A web page object that displays a live URL

B. A live SQL table object

C. A download object (lets users export the dashboard as image / PDF / PowerPoint)

D. An extension object (Tableau or third-party)

Question 49. Stakeholders want to see how each region's contribution to total monthly sales changes over 12 months - composition AND trend in one view. Which chart fits best?

A. A pie chart for each month

B. A scatter plot

C. A stacked area chart (month on the continuous axis, Region stacked on Color)

D. A Gantt chart

Question 50. On a line chart of daily revenue, only the highest and lowest points of each line should carry value labels. How do you configure this?

A. Drag the measure to Label and then manually hide the unwanted labels

B. Add a reference line labeled with the maximum value

C. Add a table calculation filter that keeps only the max and min days

D. Marks card -> Label -> enable Min/Max (allow some labels)

Question 51. A bar chart's value axis auto-scaled to 340-1,200, which visually exaggerates small differences. Management requires an honest zero baseline. What do you do?

A. Right-click the axis -> Edit Axis -> set the Range to Fixed, starting at 0

B. Add a reference line at 0

C. Convert the measure to a table calculation

D. Synchronize the axis with a secondary axis

Question 52. For a minimalist board deck, all gridlines must be removed from a crosstab. Where do you do this?

A. Worksheet menu -> Clear Sheet

B. Format menu -> Lines -> set Row and Column Gridlines to None

C. Right-click each cell -> Format Cell

D. Dashboard menu -> Actions -> remove gridline actions

Question 53. On a scatter plot of customers, each click should ADD the clicked customer to the previously selected set, growing the comparison group click by click. Which set action configuration does this?

A. Running action: Assign values to set

B. Running action: Remove from set

C. Running action: Add to set

D. Use a filter action that targets the set instead

Question 54. A filter action filters the detail sheet when a bar is selected. After the user DESELECTS the bar, the detail sheet must go back to showing everything. Which 'Clearing the selection will' option must be chosen?

A. Exclude all values

B. Leave the filtered view

C. Hide the target sheet

D. Show all values

Question 55. An executive wants a quick spatial comparison of category performance where each category is a bubble sized by revenue, with no nested hierarchy. Which chart type fits?

A. Treemap

B. Highlight table

C. Packed bubbles

D. Stacked bar

Question 56. Two teams publish conflicting 'Total Sales' data sources. Analysts must see at a glance which data source the data steward has endorsed as the official one. Which feature communicates this?

A. Certify the governed data source on Tableau Server/Cloud - a gold Certified badge appears next to it

B. Set the conflicting source's refresh schedule to run last

C. Subscribe all analysts to the governed dashboard

D. Create a custom view of the governed workbook

Question 57. A published extract has failed its last three scheduled refreshes, and analysts keep querying stale numbers. Which Server/Cloud feature warns downstream users directly, inside Tableau, that the data is not fresh?

A. A subscription to the data source

B. A data quality warning on the data source (a stale/degraded badge on the content)

C. A data-driven alert on every dashboard that uses it

D. Disabling download permissions on the workbook

Question 58. A manager right-clicks around a published text-table view looking for 'Create Alert' and cannot find it. Why?

A. Data-driven alerts can only be created by site administrators

B. Alerts require a live connection; the view uses an extract

C. Data-driven alerts attach to a view's numeric (continuous) axis - the view must plot a measure on an axis to attach a threshold

D. Alerts must be created in Tableau Desktop, never on Server/Cloud

Question 59. A user holds the Viewer site role. An admin grants her the Save capability at the project level on a workbook, but she still cannot save edits to the published workbook. Why?

A. Save permissions only apply to data sources, never workbooks

B. Workbooks are permanently locked to their original author

C. She must clear her browser cache and reload

D. The site role caps capabilities - a Viewer can interact only, so she must be upgraded to Explorer (can publish) or Creator to save content

Question 60. Content in a project has chaotic permissions inherited from a parent project. The admin wants ONE set of permissions to apply to all content and child projects from a single place. Which mechanism does this?

A. Move all content to the Default project

B. Certify the project

C. Lock the permissions to the project (Lock Permissions to Project)

D. Publish everything as .twbx files

ANSWER SHEET (for self-marking)

Record your answers here while taking the exam, then check against advanced-answers.pdf.

Q1		Q31		Q11		Q41		Q21		Q51	
Q2		Q32		Q12		Q42		Q22		Q52	
Q3		Q33		Q13		Q43		Q23		Q53	
Q4		Q34		Q14		Q44		Q24		Q54	
Q5		Q35		Q15		Q45		Q25		Q55	
Q6		Q36		Q16		Q46		Q26		Q56	
Q7		Q37		Q17		Q47		Q27		Q57	
Q8		Q38		Q18		Q48		Q28		Q58	
Q9		Q39		Q19		Q49		Q29		Q59	
Q10		Q40		Q20		Q50		Q30		Q60	

Score: ____ / 60 = ____ % | Passing the REAL exam takes 39/60 (65%). On this harder set, 39+ means you are comfortably on track; 45+ means you are exam-ready. Review every explanation in the answer key - especially for questions you guessed correctly.